

Essays on the Application of

Statistical Learning

in

Empirical Economic Research

Julian Dörr^{1,2}

¹Justus Liebig University Giessen

²former: ZEW - Leibniz Centre for European Economic Research

November 4, 2022


Outline of talk

00.1 Technology-company mapping framework


Mapping Technologies to Business Models: An Application to Clean Technologies and Entrepreneurship

published as part of the *the 26th International Conference on Science, Technology and Innovation Indicators (STI2022)* Conference Proceedings

Technological innovation and its measurement




Technological innovation and its measurement



***Patents** have become a surrogate for measuring the innovation process.*

Jaffe (2021)

Technological innovation and its measurement



***Patents** have become a surrogate for measuring the innovation process.*

Jaffe (2021)

***Patent** subclasses provide a [...] reliable picture of a firm's technological capabilities.*

Aharonson et al. (2016)

Patents and start-ups

A measurement problem

Start-ups barely file patents (Graham et al., 2008; Helmers et al., 2011):


- ▶ distracting engineers/managers from key functions (Mann, 2005)
- ▶ costs of patenting/patent litigation too high (Graham et al., 2009)
- ▶ disclosure through patent allows 'design around' (Mann, 2005)
- ▶ patents impact VC decisions as property rights, not as signals of technology quality (Hoenig et al., 2015)

Research question I

How to capture the role of start-ups in the innovation process?

Textual innovation data

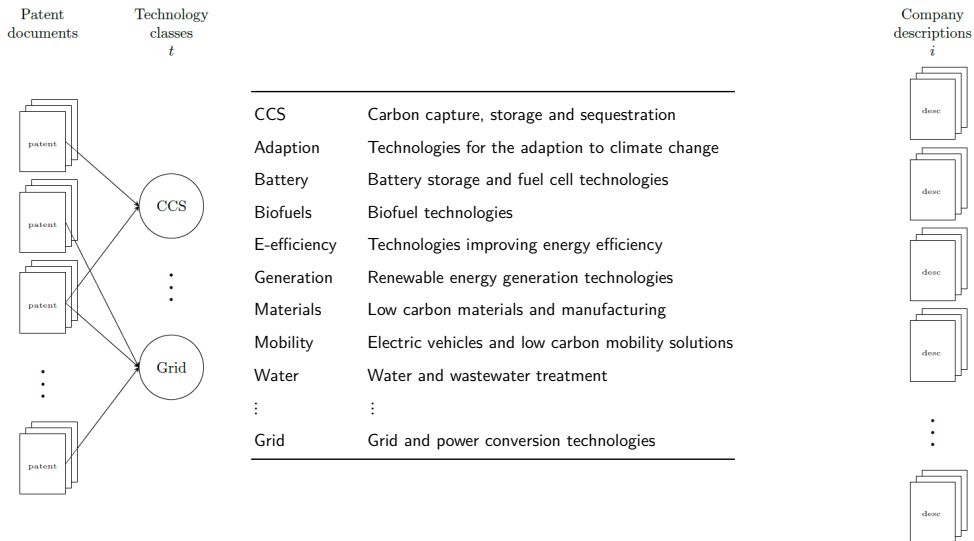
Legal obligation to publish business purpose at business registration



Source: handelsregister.de


Textual innovation data

Patent texts and assigned technology classes




From patents to technology descriptions

L-LDA (Ramage et al., 2009)

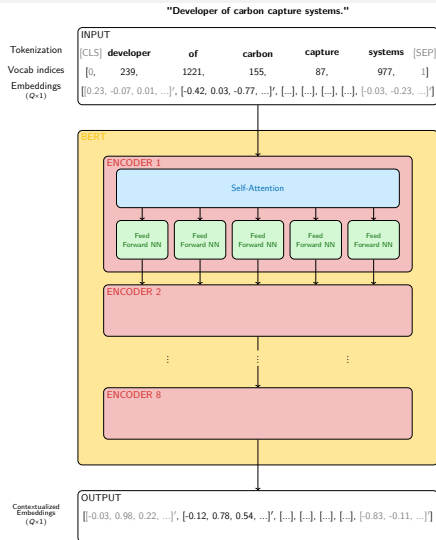


Goal: Derive technology-word distributions from expert-labeled corpus of patent docs




Excursus: BERT

Model architecture




Attention Is All You Need (Vaswani et al., 2017)

Let tokens 'look around' the whole input, and decide how to update its representation based on on what it sees



Encoder


After Attention, each token pondering for itself about what it has observed previously




Training BERT

Self-supervised learning based on English Wikipedia

1. Masked language modeling




2. Next sentence prediction



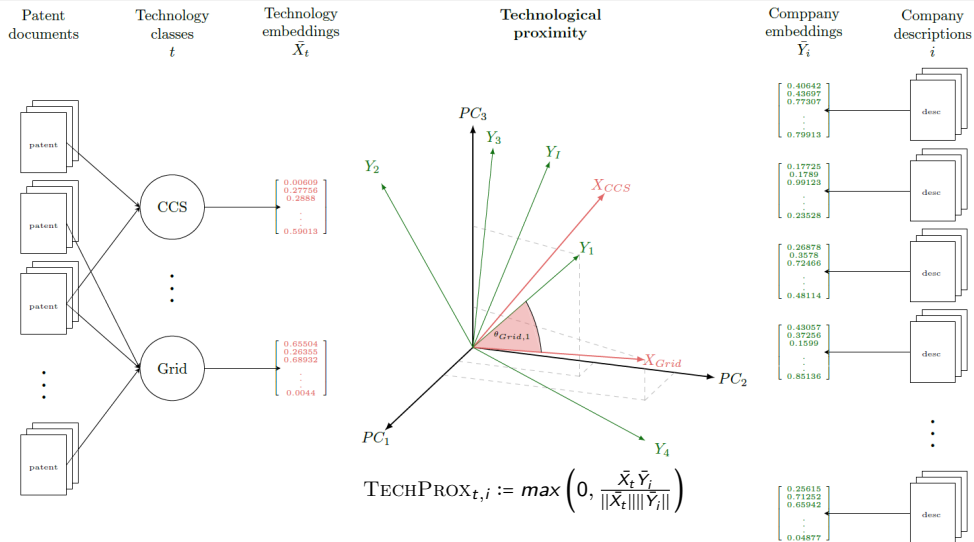
Finetuning BERT: SBERT (Reimers et al., 2019)

Finetuning based on collection of sentence pairs labeled for entailment, contradiction, and semantic independence



Mapping framework

Cosine similarity as measure of a company's technological orientation



Application

Role of start-ups in clean technology diffusion

Fight against climate change requires new technological pathways and radical innovations (*inter alia* European Commission (2019), United Nations (2015))

- ▶ but: technological path dependencies and system/innovation inertia among incumbents (Aghion et al., 2016; Patel et al., 1997)
- ▶ costly: delay in redirecting innovation towards clean technologies (Benner, 2009; Dijk et al., 2016; Sick et al., 2016)


⇒ special role of new (path-independent!) ventures in driving clean technology change

Research question II

Which role do start-ups play in the technological transition to higher levels of decarbonization?


Application

TECHPROX in survey of German start-ups




Application

A glance at the 'outliers'




Application

A glance at the 'outliers'




Application

A glance at the 'outliers'




Application

A glance at the 'outliers'



Application

A glance at the 'outliers'



Characteristics of clean technology start-ups

Cleantech start-ups show a higher propensity to eco-innovate

	<i>Elnno</i>					
	(1)	(2)	(3)	(4)	(5)	(6)
TECHPROX	1.015*	1.014*	1.013*	1.013*	1.012*	1.014***
log(size)		1.190***	1.140***	1.125***	1.186***	1.175***
age		1.001	1.010	1.001	1.005	1.012
subsidy			1.317***	1.353***	1.413***	1.456***
R&D			1.427***	1.434***	1.605***	1.675***
R&D intensity			0.780	0.910	0.904	0.815
returns				1.743***	1.633**	1.551**
break even				1.295***	1.226**	1.237**
team size					0.899**	0.887**
university					0.614***	0.627***
Sector controls	Y	Y	Y	Y	Y	Y
Product type controls	N	N	N	N	N	Y
<i>N</i>	3,269	3,269	3,269	3,192	3,192	2,774
Pseudo <i>R</i> ²	0.022	0.026	0.030	0.033	0.041	0.047

Note:

Elnno := Introduction of environmental innovation?

- no environmental innovation
- environmental innovation with moderate environmental effect
- environmental innovation with substantial environmental effect

Coefficient estimates reported as proportional odds ratios. Significance levels: *: $p < 0.10$, **: $p < 0.05$, ***: $p < 0.01$

Summary

- ▶ Latest evolutions in the field of NLP allow fine granular determination of a firm's technological profile
- ▶ Legal obligation to publish a business purpose makes the technology mapping possible for start-ups even w/o traditional innovation data
- ▶ Leveraging the introduced technology mapping to the field of clean technologies suggests:
 - ▶ a high propensity of cleantech start-ups to introduce eco-innovations
 - ▶ supporting their special role in the transition to a green economy derived from theory
 - ▶ both by virtue of their business models as well as a high propensity to adopt additional environmental innovations

- Acs, Z. J., & Audretsch, D. B. (2005). Entrepreneurship, Innovation and Technological Change. Now Publishers Inc.
- Aghion, P., Dechezleprêtre, A., Hémous, D., Martin, R., & van Reenen, J. (2016). Carbon taxes, path dependency, and directed technical change: Evidence from the auto industry. Journal of Political Economy, 124(1), 1–51.
<https://doi.org/10.1086/684581>
- Aharonson, B. S., & Schilling, M. A. (2016). Mapping the technological landscape: Measuring technology distance, technological footprints, and technology evolution. Research Policy, 45(1), 81–96.
<https://doi.org/10.1016/j.respol.2015.08.001>
- Bengio, Y., Ducharme, R., Vincent, P., & Jauvin, C. (2003). A Neural Probabilistic Language Model. Journal of Machine Learning Research, 3, 1137–1155. <https://doi.org/10.1080/1536383X.2018.1448388>
- Benner, M. J. (2009). Dynamic or static capabilities? Process management practices and response to technological change. Journal of Product Innovation Management, 26(5), 473–486.
<https://doi.org/10.1111/j.1540-5885.2009.00675.x>
- Blei, D. M., Ng, A. Y., & Jordan, M. I. (2003). Latent Dirichlet Allocation. Journal of Machine Learning Research, 3, 993–1022. <https://doi.org/10.1145/2133806.2133826>
- Bojanowski, P., Grave, E., Joulin, A., & Mikolov, T. (2017). Enriching Word Vectors with Subword Information. Transactions of the Association for Computational Linguistics, 5arXiv 1607.04606, 135–146.
https://doi.org/10.1162/tacl_a_00051
- Devlin, J., Chang, M.-W., Lee, K., & Toutanova, K. (2018). BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding. arXiv 1810.04805. <https://doi.org/10.48550/arXiv.1810.04805>
- Dijk, M., Wells, P., & Kemp, R. (2016). Will the momentum of the electric car last? Testing an hypothesis on disruptive innovation. Technological Forecasting and Social Change, 105, 77–88.
<https://doi.org/10.1016/j.techfore.2016.01.013>
- European Commission. (2019). The European Green Deal. <https://doi.org/10.2307/j.ctvd1c6zh.7>

- Firth, J. R. (1957). A synopsis of linguistic theory, 1930-1955. <http://annabellelugin.edublogs.org/files/2013/08/Firth-JR-1962-A-Synopsis-of-Linguistic-Theory-wfih5.pdf>
- Gottschalk, S. (2013). The Research Data Centre of the Centre for European Economic Research (ZEW-FDZ). *Journal of Contextual Economics – Schmollers Jahrbuch*, *133*(4), 607–618.
- Graham, S. J., Merges, R. P., Samuelson, P., & Sichelman, T. (2009). High Technology Entrepreneurs and the Patent System : Results of the 2008 Berkeley Patent Survey Author (s): Stuart J . H . Graham , Robert P . Merges , Pam Samuelson and Ted Sichelman Published by : University of California , Berkeley , School of Law S. *Berkeley Technology Law Journal*, *24*(4), 1255–1327.
- Graham, S. J., & Sichelman, T. (2008). Why Do Start-ups Patent? *Berkeley Technology Law Journal*, *23*(3), 1063–1097.
- Helmers, C., & Rogers, M. (2011). Does patenting help high-tech start-ups? *Research Policy*, *40*(7), 1016–1027. <https://doi.org/10.1016/j.respol.2011.05.003>
- Hoening, D., & Henkel, J. (2015). Quality signals? the role of patents, alliances, and team experience in venture capital financing. *Research Policy*, *44*(5), 1049–1064. <https://doi.org/10.1016/j.respol.2014.11.011>
- Howard, J., & Ruder, S. (2018). Universal language model fine-tuning for text classification. In *Proceedings of the 56th annual meeting of the association for computational linguistics*. <https://doi.org/10.48550/arXiv.1801.06146>
- Jaffe, A. B. (2021). Patent Metrics for Innovation Research. (September).
- Joulin, A., Grave, E., Bojanowski, P., & Mikolov, T. (2017). Bag of tricks for efficient text classification. *Proceedings of the 15th conference of the European chapter of the association for computational linguistics*, *2*(Short Papers), 427–431. <https://doi.org/10.1176/appi.ps.201500423>
- Mann, R. J. (2005). Texas Law Review. *Texas Law Review*, *83*(4).

- Mikolov, T., Sutskever, I., Chen, K., Corrado, G., & Dean, J. (2013). Distributed Representations of Words and Phrases and their Compositionality. Advances in neural information processing systems, arXiv 1606.08359, 3111–3119. <https://doi.org/10.18653/v1/d16-1146>
- Patel, P., & Pavitt, K. (1997). The technological competencies of the world's largest firms: Complex and path-dependent, but not much variety. Research Policy, 26(2), 141–156. [https://doi.org/10.1016/S0048-7333\(97\)00005-X](https://doi.org/10.1016/S0048-7333(97)00005-X)
- Pennington, J., Socher, R., & Manning, C. D. (2014). GloVe: Global Vectors for Word Representation Jeffrey. Proceedings of the 2014 conference on empirical methods in natural language processing (EMNLP), 1532–1543. <https://doi.org/10.1080/02688697.2017.1354122>
- Peters, M. E., Neumann, M., Iyyer, M., & Gardner, M. (2018). Deep contextualized word representations. Proceedings of the 2018 conference of the north American chapter of the association for computational linguistics: 1(Long Paper), arXiv 1802.05365, 2227–2237.
- Radford, A., Narasimhan, K., Salimans, T., & Sutskever, I. (2018). Improving Language Understanding by Generative Pre-Training. <https://doi.org/10.4310/HHA.2007.v9.n1.a16>
- Ramage, D., Hall, D., Nallapati, R., & Manning, C. D. (2009). Labeled LDA: A supervised topic model for credit attribution in multi-labeled corpora. Proceedings of the 2009 Conference on Empirical Methods in Natural Language Processing, 248–256. <http://www.aclweb.org/anthology/D09-1026>
- Reimers, N., & Gurevych, I. (2019). Sentence-BERT: Sentence embeddings using Siamese BERT-Networks. In Proceedings of the 2019 conference on empirical methods in natural language processing and the 9th international Stroudsburg, PA, USA, Association for Computational Linguistics. <https://doi.org/10.18653/v1/D19-1410>
- Sick, N., Nienaber, A. M., Liesenkötter, B., vom Stein, N., Schewe, G., & Leker, J. (2016). The legend about sailing ship effects – Is it true or false? The example of cleaner propulsion technologies diffusion in the automotive industry. Journal of Cleaner Production, 137, 405–413. <https://doi.org/10.1016/j.jclepro.2016.07.085>

- United Nations. (2015). Paris Agreement. Retrieved October 1, 2020, from https://unfccc.int/sites/default/files/english%7B%5C_%7Dparis%7B%5C_%7Dagreement.pdf
- Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A. N., Kaiser, L., & Polosukhin, I. (2017). Attention Is All You Need. 31st Conference on Neural Information Processing Systems (NIPS 2017), arXiv 1706.03762. <http://arxiv.org/abs/1706.03762>
- Wang, Y., Hou, Y., Che, W., & Liu, T. (2020). From static to dynamic word representations: a survey. International Journal of Machine Learning and Cybernetics, 11(7), 1611–1630. <https://doi.org/10.1007/s13042-020-01069-8>

Appendix

Technology-company mapping framework

"Developer of carbon capture systems."

Tokenization
Vocab indices
Embeddings
($Q \times 1$)

INPUT

[CLS] **developer** **of** **carbon** **capture** **systems** [SEP]

[0, 239, 1221, 155, 87, 977, 1]

[[0.23, -0.07, 0.01, ...]', [-0.42, 0.03, -0.77, ...]', [...], [...], [...], [...], [-0.03, -0.23, ...]']

BERT

ENCODER 1

BERT

ENCODER 1

Self-Attention

Feed
Forward NN

Feed
Forward NN

Feed
Forward NN

Feed
Forward NN

Feed
Forward NN


ENCODER 2

⋮

⋮

⋮

ENCODER 8



ENCODER 6

Contextualized
Embeddings
($Q \times 1$)

OUTPUT

$[[-0.03, 0.98, 0.22, \dots]', [-0.12, 0.78, 0.54, \dots]', [\dots], [\dots], [\dots], [\dots], [-0.83, -0.11, \dots]']$

w_1 **developer**

[-0.42, 0.03, -0.77, ...]'

 w_2 **of**

[...]

 w_3 **carbon**


[...]

 w_4 **capture**

[...]


 w_5 **systems**



[...]



SELF-ATTENTION

SELF-ATTENTION





1. Attention weights $a_{1:5}$ are query-key similarities:

$$\hat{a}_i = \mathbf{q}_i \times \mathbf{k}_i$$

Normalized via softmax: $a_i = e^{\hat{a}_i} / \sum_j e^{\hat{a}_j} \in [0, 1]$


2. Output $\hat{\mathbf{z}}_i$ is attention-weighted average of value vectors $\mathbf{v}_{1:5}$:
(1×Z)

$$\hat{\mathbf{z}}_i = \sum_j a_j \mathbf{v}_j$$


3. \mathbf{k} , \mathbf{v} and \mathbf{q} are derived from the entire input \mathbf{w} :

$$\mathbf{k} = W_k \times \mathbf{w} \quad \mathbf{v} = W_v \times \mathbf{w} \quad \mathbf{q} = W_q \times \mathbf{w}$$


Note: Self-attention is repeated H times (multi-head attention) and the resulting vectors are concatenated along the feature dimension. Multiplying with a weight matrix W_z yields the final output vector that is passed to the FNN.
(Q×HZ)






1. Masked language modeling



2. Next sentence prediction






contra-
diction

entailment

neutral



A patent reflects new technical knowledge, but it does not indicate whether this knowledge has a positive economic value. Only those inventions which have been successfully introduced in the market can claim that they are innovations as well. While innovations and inventions are related, they are not identical.

Acs et al. (2005)

Text preprocessing

1. translation of non-English texts to English
2. Part of Speech (PoS) tagging
 - 2.1 remove punctuation, numbers and unknown tags
 - 2.2 lemmatization
3. stop word deletion

A labeled corpus of patent abstracts

Patent	Technology class	Abstract
1	B, C, Y02C, Y02P	Catalyst, comprising one or more compounds of the perovskite-type as catalytically active component, is new, where the catalytically active component in the form of at least one layer is applied on a support body from an open cell foam ceramic material ...
2	A, Y02A, Y02C, Y02E	Absorber fluid, comprises a carbon dioxide binding absorbent and an ionic additive in a concentration, which is greater than a minimum concentration, so that the activity of the products formed by the connection of carbon dioxide to the absorbent is reduced ...
⋮	⋮	⋮
P	B, F, Y02C	The invention relates to a power plant for generating electrical energy, comprising a combustion chamber for producing steam, at least one waste gas purification stage that is connected downstream, a separation stage for CO ₂ ...

Note: Corpus comprises $P \sim 560,000$ patents (all patents filed by German firms after 1990) and a vocabulary size of $V \sim 370,000$ (after `text preprocessing`).

Vertical differentiation in technology classes

Classification system of the European Patent Office using the example of **carbon capture and storage technologies**:

CPC	COOPERATIVE PATENT CLASSIFICATION
Y	New technological developments
Y02	Climate change mitigation technologies
Y02C	Carbon capture and storage technologies
Y02C20	Capture and disposal of greenhouse gases
Y02C20/10	- of N_2O

Latent Dirichlet Allocation

Core idea in Blei et al. (2003) seminal work on Latent Dirichlet Allocation (LDA):

Model the generative process that led to the creation of a text corpus incorporating both:

- ▶ the observed words in the corpus' documents
- ▶ *and* the hidden topic structure within the corpus

in the imaginary data generating process.

The latter includes the distribution of topics over documents and the word distributions over topics.

Latent Dirichlet Allocation

Core idea in Blei et al. (2003) seminal work on Latent Dirichlet Allocation (LDA):

Model the generative process that led to the creation of a text corpus incorporating both:

- ▶ the observed words in the corpus' documents
- ▶ *and* the hidden topic structure within the corpus

in the imaginary data generating process.

The latter includes the distribution of topics over documents and the word distributions over topics.

L-LDA (Ramage et al., 2009) extends upon LDA by taking into consideration document labels in the generative process.

L-LDA in patent corpus:

- ▶ document $\hat{=}$ patent, p
- ▶ labels/topics $\hat{=}$ technology classes, t
- ▶ word distributions over topics $\hat{=}$ semantic technology description, $p(\delta_t)$

Statistical Learning in L-LDA

Patent corpus D consisting of P distinct patent abstracts each of length N_p , generative process can be modeled as follows:

1. For each technology class $t \in \{1, \dots, T\}$: generate word distribution $\delta_t \sim \text{Dir}(\beta)$
2. For each patent $p \in \{1, \dots, P\}$: generate technology class distribution $\lambda_p \sim \text{Dir}(\alpha_p)$
3. For each of the word positions p, n , with $p \in \{1, \dots, P\}$ and $n \in \{1, \dots, N_p\}$:
 - 3.1 generate technology class assignment $z_{p,n} \sim \text{Multinomial}(\lambda_p)$
 - 3.2 and choose word $w_{p,n} \sim \text{Multinomial}(\delta_{z_{p,n}})$

$$p(\delta_{1:T}, \lambda_{1:P}, z_{1:P}, w_{1:P}) = \prod_{t=1}^T p(\delta_t) \prod_{p=1}^P p(\lambda_p) \left(\prod_{n=1}^{N_p} p(z_{p,n} | \lambda_p) p(w_{p,n} | \delta_{z_{p,n}}) \right)$$

Goal: Derive posterior distribution $p(\delta_t)$ from joint distribution $p(\delta_{1:T}, \lambda_{1:P}, z_{1:P}, w_{1:P})$

Importance of capture contextual meaning of words

- ▶ **technical terms** in technology descriptions:

$X_t = \langle \text{gas, absorb, carbon, dioxide, desorption} \dots \rangle$

- ▶ **non-technical terms** in company descriptions:

'Developer of direct air capture technology that safely and permanently removes CO2 from the air.'

$\rightarrow Y_c = (\text{developer, direct, air, technology, safe, permanent, remove, co2})'$

- ▶ **But:** high semantic overlap between x_t and y_c as captured by token embeddings

$\bar{X}_t(\text{carbon}) \approx \bar{Y}_c(\text{co2})$

$\bar{X}_t(\text{absorb}) \approx \bar{Y}_c(\text{remove})$

- ▶ **Goal:** Exploit these relations to capture adopters of a technology


Classification performance of TECHPROX

Table: Performance of TECHPROX in distinguishing cleantech from non-cleantech firms

Label	Precision	Recall	F1-Score	Support
Cleantech	0.87	0.86	0.86	284
Non-cleantech	0.83	0.84	0.83	233
			0.85	517

Note: Performance measured on random test set with optimal values of $Q = 15$ and $\text{TECHPROX}_{\min} = 0.27$. Optimal values for Q and TECHPROX_{\min} have been determined on the validation set by tuning F1-Score.

Evolution of NLP



Word embeddings (1)

You shall know a word by the company it keeps!

Firth (1957)

General idea: exploit information on co-occurrence of words in large text corpora in order to learn the semantic meaning of a word as represented by a low-dimensional, dense vectors ($E \ll V$).

Natural Language Processing (NLP) as highly active field of research with major advances in recent years (see Wang et al. (2020)):

Neural Network Language Models

- ▶ 'distributed representation for words' (Bengio et al., 2003)
 - ▶ learn model that predicts next word given previous words
 - ▶ word embeddings carrying semantic meaning of a word as by-product

Word embeddings (2)

Static word embeddings

- ▶ Word2Vec (Mikolov et al., 2013)
 - ▶ neural network architecture specifically designed to learn word embeddings
 - ▶ Continuous Bag-of-Words (CBOW): predict word given its surrounding context words
 - ▶ Skipgram: predict context words given central word
- ▶ GloVe (Pennington et al., 2014)
 - ▶ direct exploitation of co-occurrence statistics from large text corpora
- ▶ fastText (Bojanowski et al., 2017; Joulin et al., 2017)
 - ▶ learning embeddings for character n-grams and representing words as the sum of the n-gram embeddings (towards multi-language models)

Word embeddings (3)


Contextualized word embeddings

Tackle the issue that words have different meanings in different contexts (polysemy)

- ▶ ELMo (Peters et al., 2018)
 - ▶ use bidirectional LSTM to capture whole sentence (context!) in order to model embeddings of words in sentence
- ▶ ULMFit (Howard et al., 2018)
 - ▶ introduce a general language model and a process to fine-tune to domain-specific NLP tasks
- ▶ GPT (Radford et al., 2018)
 - ▶ use transformer network architecture to learn linguistic long-term dependencies
- ▶ BERT (Devlin et al., 2018)
 - ▶ Consider bidirectional contexts and relation of sentence pairs based on transformer encoders


Application

A glance at the 'outliers'



Application


A glance at the 'outliers'



Application


A glance at the 'outliers'

Manufacture of electrode foils, lithium accumulators and energy storage systems and the provision of services in this area.




Application

A glance at the 'outliers'



Application

A glance at the 'outliers'



IAB/ZEW Start-up survey

A representative sample of German start-up companies (Gottschalk, 2013)

Table: 2018 IAB/ZEW Start-up survey questions on environmental impacts and environmental innovation

Environmental impact

Does your company offer products or services which have the following environmental effects on the customer or the end user?

1. Reduction of energy consumption or CO₂ footprint for the customer.
2. Reduction of other emissions to the air, water, soil or noise for the the customer.
3. Reduction of material or resource consumption, for instance water, for the customer.
4. Improvement of recyclability of customer's products.
5. Improvement of durability of customer's products.

Environmental innovation

Since its inception, has your company introduced innovations that have impacted the environment as follows?

1. Reduction of energy consumption or the overall CO₂ balance in your company.
2. Reduction of other emissions to the air, water, soil or noise in your company.
3. Reduction of material or resource consumption, for instance water, in your company.
4. Improvement of recyclability of your own products.
5. Improvement of durability of your own products.

Note: The questions have been asked on a Likert response scale with the following response possibilities. (1) No; (2) Yes, somewhat; (3) Yes, substantial.